Thornwood Rescue: Debrief

1. Kaito: Why did I get stuck on the first hill?

Where it stopped: (2,1) with value **70**, after 3 steps (`HC_FINAL`, `HC_VALUE`, `HC_STEPS`). The climb went (0,0)=52 → (0,1)=58 → (1,1)=64 → (2,1)=70. At (0,0) the right and down neighbours tie at 58, and the up, right, down, left rule takes right.

around (2,1)	col 0	col 1	col 2
row 1	58	64	58
row 2	64	**70**	64
row 3	58	64	58

Its neighbours: up (1,1)=64, right (2,2)=64, down (3,1)=64, left (2,0)=64. None is strictly higher than 70, so the stopping rule fires.

Why this is a local maximum: Hill climbing only compares a cell with its four neighbours, so any cell that beats all four looks like the top. (2,1) is one of only two such peaks on the map; the other is the real summit (6,6)=95. The best route between them still steps on a 53, a drop of 17, and steepest ascent never goes down. Restarting is the cure: the climbs from (4,0), (7,0) and (0,4) end on 70, 95 and 95 after 3, 7 and 8 steps, so the best of the four is 95 from (7,0), 21 steps in total (`RR_BEST_VALUE`, `RR_BEST_START`, `RR_TOTAL_STEPS`).

2. Gennai: Why did only some of the 20 fate scrolls reach the summit?

- **Result of the 20 seeded runs** (T0 = 20, cooling 0.97): **8 of 20** end on 95 (`SA_SUCCESS_1_TO_20`). The other 12 all freeze on the same false summit, (2,1)=70, the one that trapped Kaito.
- **How heat and cooling act.** A downhill step $\Delta < 0$ is accepted with chance $\exp(\Delta/T)$, and T becomes $0.97 \cdot T$ every iteration, so $T = T0 \cdot 0.97^n$. Leaving the false summit means a drop of 17 (Q1), so the first table gives the chance of accepting that step as the run cools.

Iteration n	T (T0=20)	chance of −17	T (T0=40)	chance of −17
0	20.00	0.427	40.00	0.654
25	9.34	0.162	18.68	0.402
50	4.36	0.020	8.72	0.142
100	0.95	0.000	1.90	0.000
150	0.21	0.000	0.41	0.000
199	0.05	0.000	0.09	0.000

successes of 20	rate 0.95	rate 0.97	rate 0.99
T0 = 10	1	3	11
T0 = 20	3	8	10

successes of 20	rate 0.95	rate 0.97	rate 0.99
T0 = 40	8	12	7
T0 = 80	11	12	2

- Early on the walker is bold; once T < 1 (after iteration 98 for T0=20) it is greedy. Accepted downhill moves per quarter of the run, summed over the 20 seeds: 269, 61, 0, 0 (T0=20) and 336, 169, 6, 0 (T0=40). Once the chance of crossing falls below 1% (after iteration 55 for T0=20, 78 for T0=40) a run still on the 70-hill is frozen there.
- Cooling rate works the same way (second table above). With T0 = 20, fast cooling shuts the window early (3 of 20 at 0.95) and slow cooling keeps it open (10 of 20 at 0.99). More heat is not always better: T0 = 80 with 0.99 gives only 2 of 20, because T is still 10.7 after 200 iterations and the walker is still wandering when the run stops. A good run is hot enough to cross the dip early, then cold enough to settle on the peak.
- Re-run with T0 = 40: 12 of 20 reach 95 (was 8). The 8 failures again all sit on (2,1)=70. The crossing window stays open for 78 iterations instead of 55, so more runs cross before freezing and still have about 122 iterations to climb and settle.

3. Rin: Why did the seeds stop at 234 when 241 was possible?

The GA pouch is `100111100110001` and the optimum is `101110100100101`. The four tools where they disagree are in bold.

Tool	Load	Vol	Use	GA	Opt
0 Kunai set	18	9	36	1	1
1 Wire spool	11	8	9	0	0
2 Smoke bombs	**12**	**5**	**10**	**0**	**1**
3 Paper bombs	9	3	32	1	1
4 Energy pills	5	7	31	1	1
5 Flash tags	**5**	**6**	**8**	**1**	**0**
6 Messenger scroll	11	3	33	1	1
7 Giant shuriken	20	6	27	0	0
8 Spare rations	13	6	5	0	0
9 Wind shuriken	11	2	35	1	1
10 Twin blades	**14**	**6**	**21**	**1**	**0**
11 Caltrops	15	9	13	0	0
12 Medical kit	**8**	**8**	**26**	**0**	**1**
13 Grappling hook	10	9	19	0	0
14 Binding tags	10	7	38	1	1
GA pouch total	83 / 86	43 / 44	**234**		

Tool	Load	Vol	Use	GA	Opt
Optimal pouch total	84 / 86	44 / 44	**241**		

- Differences are that the GA packs Flash tags (use 8) and Twin blades (use 21); the optimum packs Smoke bombs (use 10) and Medical kit (use 26) instead. That is +7 usefulness (29 → 36) for +1 load and +1 volume. The GA pouch holds 83 of 86 load and 43 of 44 volume; the optimum uses 84 and exactly 44.
- The seeds missed since the pouches are 4 bit flips apart and no better pouch exists within 3 flips of the GA's (better pouches within 1, 2, 3, 4 flips: 0, 0, 0, 1). Of the 14 half-way pouches, 8 break a limit and score 0 (two, worth 252 and 239, miss by only 1 volume unit) and 6 are legal but score below 234. So the GA must make all four changes at once, and by then diversity had run out: 20 distinct pouches at generation 0, 19 at generation 15 (when 234 first appeared and the best never changed again), 14 at generation 50, with 6 of the 20 breaking a limit (hence the average of 148.20).
- One fix could be to raise the mutation rate from 1/15 to 2/15 per bit. On seed 12 the GA then reaches 241 at generation 24, while at 1/15 it stays on 234 even if run for 150 generations. Honest limits: over seeds 1 to 200 the hit rate only moves from 87 to 91 of 200, and 3/15 is worse (57 of 200). Longer runs lift it to 141 of 200 but do not rescue seed 12.

4. Vesper: How would I have beaten you on every losing move?

Vesper's diagonal (0,4), (1,3), (3,1) is one stone short, and the empty cell (2,2) is where a stone dropped in column 2 lands. After any first move except column 2 she plays column 2 and wins on her very next stone. Only column 2 blocks it.

X's first move	Value	Vesper's reply	Result
0	−10	O2 completes (0,4), (1,3), (2,2), (3,1)	Vesper wins at once, 9 cells left: −(1+9) = −10
1	−10	O2 completes (0,4), (1,3), (2,2), (3,1)	Vesper wins at once, 9 cells left: −(1+9) = −10
2	+1	none wins; best defence O0 X0 O1 X2 O2 X3 O4 X4 O4 X4	X completes (1,1), (2,2), (3,3), (4,4) on a full board: +1
3	−10	O2 completes (0,4), (1,3), (2,2), (3,1)	Vesper wins at once, 9 cells left: −(1+9) = −10
4	−10	O2 completes (0,4), (1,3), (2,2), (3,1)	Vesper wins at once, 9 cells left: −(1+9) = −10

5. Jun: Compare the ki (node counts)

Method	Nodes	Share
Full minimax (0..4)	175,789	baseline
AB1: alpha-beta, full, 0..4	9,064	5.2% of minimax
AB2: alpha-beta, full, middle first	8,965	5.1% of minimax
Depth-limited minimax (depth 4)	421	baseline

Method	Nodes	Share
AB3: alpha-beta, depth 4	198	47.0% of depth-limited

The pruned runs return the same values as the plain searches: AB1 and AB2 give 1 like full minimax, and AB3 gives −2 like the depth-limited search. Pruning only skips work; it never changes the answer.

- **Why middle-first barely helped.** It saved 99 nodes, 1.1% of what AB1 already needed. Alpha-beta only gains when the first move tried at a node is its best move (or, for the opponent, its refutation), and a fixed column order is a weak guess at that: the best column depends on the position, not on being central. I tried all 120 fixed orders: they range from 2,284 to 29,284 nodes (median 8,188), and 0..4 and middle-first rank 66th and 65th of 120, right in the middle of the pack. Trying the truly best move first at every node would need only 795 nodes; the worst-first order needs 82,690.
- **Master Hyo's tree, left to right** (hand trace below): value 9 through branch B, 39 leaves read; node B2 is called with window (−15, 26). The root opens A with (−∞, ∞). A returns −15, so the root's α becomes −15. B is searched with α = −15, returns 9, and α rises to 9. C then meets α = 9: its first child C1 returns 5, which is ≤ α, so C2 is never read.

Node	Window in (α, β)	Returns	Leaves	Cut
R (MAX)	(−∞, ∞)	9	39 / 60	–
A (MIN)	(−∞, ∞)	−15	20 / 29	–
A1 (MAX)	(−∞, ∞)	−15	10 / 12	–
A2 (MAX)	(−∞, −15)	0	7 / 7	–
A3 (MAX)	(−∞, −15)	−15	3 / 10	skips A3b, A3c, A3d
B (MIN)	(−15, ∞)	9	16 / 19	–
B1 (MAX)	(−15, ∞)	26	11 / 11	–
B2 (MAX)	(−15, 26)	9	5 / 8	–
C (MIN)	(9, ∞)	5	3 / 12	skips C2
C1 (MAX)	(9, ∞)	5	3 / 7	–

Leaves read	L→R	R→L
A	20 / 29	5 / 29
B	16 / 19	8 / 19
C	3 / 12	10 / 12
all	**39** / 60	**23** / 60

Why right to left reads fewer (23 against 39). The order decides which window each sibling is searched with. Left to right spends 20 leaves on A, the biggest and worst subtree (29 leaves, worth −15), before the root has any α to compare with. Right to left meets C, then B, and reaches A last with α = 9: A's first child returns a value ≤ α, so A2 and A1 are cut and A costs 5 leaves. Inside B it also tries B2 (9) before B1, so B1 is cut after one child. C is the only subtree that gets dearer (3 → 10). Net saving: 16 leaves.